

1 Artificial Intelligence

1.1 Introduction to Artificial Intelligence

- Artificial Intelligence is making machines that can...
 - Think like people (philosophical approach)
 - Think rationally (pass the turing test)
 - Act like people (mimic human behaviour)
 - Act rationally (the focus of our course)

- To be rational is to maximally achieve pre-defined goals. A better title for this course would be computational rationality, but that sounds less interesting.

Remark 1. *Maximize your expected utility*

- An agent is an entity that perceives and acts, our rational agents will try to maximize a utility function.

1.2 Uniformed Search

- A *reflex agent* is one that chooses it's actions based on it's perceptions and maybe some memory, it will not consider the future consequences of it's actions instead *consider how the world IS*.

Remark 2. *In our world, we will punish our agents over time to force it being optimal.*

- When developing planning agents, we need to consider:

- Ask "what if"
- Make decisions based on the predicted consequences of our actions
- Must have a model of how the world evolved to respond to actions
- Must formulate a goal(test)
- Consider how the world WILL BE

- Additionally there are different approaches to planning we can take,

- Optimal Planning (Achieve the goal in minimum cost)
- Complete Planning (When there exists a solution it is found)

- A **search problem** consists of:

- A state space (Set of possible configurations of the board)
- A successor function (How the world works, possible actions/cost of action)
- A start state and a goal state

- A solution is a sequence of actions(a plan) which transforms the state space to a goal state

- Example Problem: Traveling in NYC

- State space: Subway Stations
- Successor Function: Going to the adjacent station with cost = Time
- Start State: 34th Street Herald Square
- Goal Test: JFK Airport

- A **world state** includes every last detail about the environment

- A **search state** will only keep the details needed for planning (abstraction)

- Example Problem: Pacman Eat-All-Dots

- Heuristic: Estimate of the distance to nearest goal for each state
- A common case: Best-First takes you straight to the (wrong) goal

- What if we could combine the features of UCS and Greedy search?

- Uniform-Cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- A^* Search orders by the sum $f(n) = g(n) + h(n)$

Remark 3. *We should not stop when we enqueue a goal, but only once we pop the goal off the queue*

- A^* is sometimes optimal, but this relies on us choosing a good heuristic.

Definition 1. *A heuristic h is **admissible** (optimistic) if*

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal. An example of a admissible heuristic is the $L1$ Norm.

- Optimality of A^* Proof

- Assume:

- A is an optimal goal node, B is a suboptimal goal node, h is admissible
- Claim: A will exit the fringe before

- Proof:

- If B is never on the fringe, Q.E.D.
- Imagine if B is on the fringe
- We know some ancestor n of A is on the fringe, too (possibly A). This is because the only way for A's ancestors to be gone is if A was already expanded.
- Claim: n will be expanded before B
- 1. $f(n) \leq f(A)$
 - $f(n) = g(n) + h(n)$
 - $f(n) \leq g(A)$ by admissibility
 - $g(A) = f(A)$ since $h=0$ at goal
- 2. $f(A)$ is less than $f(B)$
 - $g(A) < g(B)$ since B is suboptimal
 - $f(A) < f(B)$ since $h=0$ at goal
- 3. n expands before B
 - $f(n) \leq f(A) < f(B)$
- Repeat for all ancestors of A until A itself is on the fringe therefore A^* search is optimal

- Uniform cost will equally fast in all directions, meanwhile A^* will tunnel directly towards a goal in an environment


- Properties of the different search algorithms

- BFS will find a path in the least amount of moves expanding outwards in each direction.
- Greedy will find the thing closest to the goal based on the heuristic, does not care about cost.
- UCS will explore forward quickly in regions where cost is lower, but does now orient itself towards the goal.
- DFS will explore the whole search space.
- A^* will explore both the heuristic and in the minimal cost direction.

- States: ((x,y), dot flags) location + boolean if dot was eaten

- Actions: UDLR
- Successor: Update location and one of the dot flags
- Goal Test: Dot Flags all flipped sign

- From these search problems we can make a **State Space Graph/Trees** or a mathematical representation of a search problem where the nodes are abstracted world configurations.



```
function Tree-Search(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- There's a problem with this strategy, we need some way to picking the fringe candidates for expansion.

- One such approach would be a **Depth-First Search**

- Strategy: Expand the deepest node first
- Implementation: Fringe is a LIFO Stack

- For the properties of the DFS:

- Let each node have factor b (each node spawns b children)
- Let m be the maximum depth of the tree
- Starting from the root node, we have one node on layer 1. On layer 2, the one root node will spawn b children.
- On layer 3, the b nodes will each have another b children therefore there will be b^2 nodes. At the end, we can state that there will be b^m nodes in the last layer.
- Time complexity will be $O(b^m)$ if m is finite.
- Each element on the fringe will be m deep in the tree, thus it will take up m space.
- Space complexity will be $O(b * m)$
- Assuming the tree cannot be infinite, then it will always find a solution.
- Does not find the optimal solution.

- Another approach we can use is **Breadth-First Search**

- As heuristics get closer to the true cost, the fewer nodes will be expanded but the more work per node to compute the heuristic will be done

- Graph Search**

- Idea: Never expand a state twice
- Tree Search + set of expanded states (the closed set)
- Only expand a node if it has never been expanded before
- If it has been expanded before, skip the expansion and add it to the closed set
- Important: store the closed set as a set and not a list
- Graph search is still complete

- Issue: If we explore an optimal node that we already found unoptimally, than the graph search will not explore it even if it would lead to a better result.

- Main Idea: Estimated Heuristics \leq actual costs

- Admissibility: heuristic cost \leq actual cost to goal
- Consistency:** heuristic "arc" cost \leq actual cost for each arc. This statement is stronger than Admissibility.
- Consequences: The $f(A)$ value along the path never decreases

- Consistency implies Admissibility

- A^* uses both backwards and (estimates) forwards costs

- A^* is optimal with an admissible/consistent heuristic

- Heuristic design is key: often use relaxed problems

1.4 Constraint Satisfaction Problems

- Standard Search Problems

- State is a black box with little knowledge of what's inside it
- Goal test can be any function over states
- Successor Function can also be anything

- Constraint Satisfaction Problems (CSPs)**

- A special subset of search problems
- State is defined by variables X, with values from a domain D.
- Goal test is a set of constraints specifying allowable combinations of values for subsets of variables.

- CSPs are a Simple example of a formal representation language

- A solution to a CSP will be one where all the assignments satisfy all the constraints

- Arcs can be used to make Constraint graphs, however they do not tell information about the constraints themselves. Example: Sudoku

- Variables: Each open square
- Domains: {1,2,3,4,5,6,7,8,9}
- Constraints: alldiff(columns), alldiff(row), alldiff(region)

- Discrete Variables

- Finite Domains

- Strategy: Expand the shallowest node first

- Implementation: Fringe is a FIFO Queue

- For the properties of the BFS:

- Let each node have factor b (each node spawns b children)
- Let depth of the shallowest solution be s
- Time complexity will be $O(b^s)$
- BFS will expand the most amount of nodes onto the fringe, as it expands each node on every layer.
- Space complexity will be $O(b^s)$
- Assuming the tree cannot be infinite, then it will always find a solution.
- Does not find the optimal solution.

- BFS outperforms DFS if the goal is shallow or on the left side of the tree

- DFS outperforms BFS memorywise

- What if we want the DFS's space advantage and the BFS's time/shallow solution advantages? Introducing **Iterative Deepening**

- Run a DFS with depth limit 1. If there's no solution...
- Run a DFS with depth limit 2. If there's no solution...
- Run a DFS with depth limit 3. ...

- While this method is redundant, generally most of the work happens in the lowest level so it isn't nearly as bad.

- When cost is present, we can use **Uniform Cost Search**

- Strategy: Expand the cheapest node first
- Implementation: Fringe is a Priority Queue (Priority: cumulative cost)

For the properties of UCS:

- Let the optimal solution has the cost C^*
- Let each individual action cost ϵ
- We can say that the effective depth of the solution is C^*/ϵ actions to get there, Furthermore the size of the fringe will also be C^*/ϵ .
- Time complexity will be $O(b^{C^*/\epsilon})$
- Space complexity will be $O(b^{C^*/\epsilon})$
- Assuming the tree cannot be infinite, then it will always find a solution.
- Will find the optimal solution! Everything on the fringe will have a higher cost than it.

- All these searches explore in every direction and doesn't think about where the goal will be.

1.3 A^* and Heuristics

- A **heuristic** is a function that estimates how close a state is to a goal. Some examples include the Manhattan Distance and the Euclidean Distance(L1 and L2 norms).

- Some characteristics of Heuristics

- These are measures, we want $h(x) = 0$ when we are at the goal state (If $y = \text{goal}$, $h(x) = |x - y| = 0$)

- Greedy Search

- Strategy: Expand the node that seems the closest to the goal state

- Size of d means $O(d^n)$ complete assignments
- Infinite domains (int, str, etc)
- Linear constraints solvable, nonlinear decidable

- Continious Variables

- Linear constraints solvable in polynomial time by LP methods

- Constraint Types

- Unary Constraints (single variable, $x \neq 0$)
- Binary Constraint (pair of variables, $x \neq y$)
- High Order Constraints involve 3 or more variables

- Standard Search Formulation

- Standard Search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
 - Initial State: the empty assignment
 - Successor function: assign a value to an unassigned variable
 - Goal test: the current assignment is complete and satisfies all the constraints

- Backtracking Search

- Variable assignments are commutative, so fix ordering($x = 1$ then $y = 2$ is the same as $y = 2$ then $x = 1$)
- Only consider assignments to a single variable at each step
- Check constraints as you go
- Incremental goal test
- Might have to do some computation to check the constraints

```
function Backtracking-Search(csp) returns solution/failure
  return Recursive-Backtracking({}, csp)
function Recursive-Backtracking(assignment, csp) returns soain/failure
  if assignment is complete then return assignment
  var = Select-Unassigned-Variable(Variables[csp], assignment, csp)
  for each value in Order-Domain-Values(var, assignment, csp) do
    if value is consistent with assignment given Constraints[csp] then
      add {var = value} to assignment
      result = Recursive-Backtracking(assignment, csp)
      if result != failure then return result
      remove {var = value} from assignment
  return failure
```

- An arc $X \rightarrow Y$ is consistent $\iff \forall x$ in the $\text{dom } X$ in the head which could be assigned without violating a constraint

- On each assignment if we enforce consistency then we can make better filtering algorithms (This method is called Arc-Consistency)

- Variable Ordering: Minimum Remaining Values(MRV) - Pick the variable with the fewest legal left domain values in its domain. It's also called fail fast ordering.

1.5 Searching with Minimax

- We can classify games as (deterministic or stochastic), (player count), (zero sum), (perfect information)
- We want algorithms that calculate a strategy(policy) which recommends a move from each state.
- In this lecture we will focus on **Deterministic games**
 - States: $S(\text{start at } s_0)$
 - Players: $P=\{1,...,N\}$ (usually take turns)
 - Actions: A (may depend on player/state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test: $S \rightarrow \{t, f\}$
 - Terminal Utilities: $S \times P \xrightarrow{D}$

- Solution for a player is a **policy**: $S \rightarrow A$

• Zero-Sum Games

- Agents have opposite utilities (values on outcomes)
- Let us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition

• General Games

- Agents have independant utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible

- To solve Zero-Sum games, we need to think about what opportunities the opponent will make, then what

- Single Agent Trees are ones where a single agent wants to achieve the best achievable outcome(utility) from that state.

- On an Adversarial Game Tree, the player and the enemy take turns doing actions. At the end of the game, we assign a utility.

- States under Agent's Control
 $V(s) = \max_{s' \in \text{successor}(s)} V(s')$
- States under Enemy Control
 $V(s') = \min_{s'' \in \text{successor}(s')} V(s'')$

- If we have two players playing against each other optimally, we know that the game will result in a draw in the case of Tic-Tac-Toe and Checkers.

- **Minimax** is when we assume the enemy will also be playing optimally.

```
def max-value(state):
    initialize v = -inf
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    initialize v = +inf
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```




Figure 2: Expectimax

• Expectimax Search

- Maybe we don't know everything about the game, sometimes there is explicit randomness
- In this case we should model it with a probability distribution

- Values should now reflect the average case(expectimax) outcome, not worst-case(minimax) outcomes

- When we run expectimax search we should compute the average score under optimal play

- The max nodes work the same as usual, but the chance nodes work by taking the weighted average of the outcomes

• Expectimax Pseudocode

```
def max-value(state):
    initialize v = -inf
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```

- A random variable represents an event whose outcome is unknown

- A probability distribution is an assignment of weights to outcomes

- Dangerous Optimism - Assume chance when the world is adversarial
- Dangerous Pessimism - Assume the worst case when it's not likely


- In Expectiminimax the environment is an extra "random agent" player that moves after each min/max nodes

Definition 2. Principle of maximum expected utility: A rational agent should choose the action that maximizes its expected utility, given its knowledge

- An agent must have preferences among:

- This is a recursion where we jump from min to max, but we will need a base case.

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is MIN: return min-value(state)
```



- Assuming root = level 1 = player's turn = max, level 2 = enemy turn = min, and level 3 = utility we can state

- $A1 = \min(3, 12, 8, \infty) = 3$
- $A2 = \min(2, 4, 6, \infty) = 2$
- $A3 = \min(14, 5, 2, \infty) = 2$
- $\text{root} = \max(3, 2, -\infty) = 3$


- Minimax is good only if both players are playing optimally, but in cases when they aren't it might be worth taking risks to gain an even bigger rewards

- Properties of Minimax:

- Minimax is like an exhaustive DFS
- Time Complexity is $O(b^m)$
- Space Complexity is $O(bm)$

- However we don't always need to look at the full tree, we can prune some of the tree to reduce the search space drastically.

- Let's walk through the Minimax Pruning of this tree



- We solve out the A1 branch to determine that our current max is 3
- We take a look inside A2 and find a 2, since we 2 is less than 3, we would never pick this option since it is less than our currently assumed max
- Inside A3, we see a 14 and that's fine. We continue exploring and see a 5, that's still better than 3 so we continue exploring. Lastly we see a 2 and compute the min between the 3 values and infinity.
- Inside this tree we skip A22 and A23 but in bigger trees there is more to save

- This technique specifically is called **Alpha-Beta Pruning**, this is an example configuration for the MIN variant

- We compute the Min value at some node n
- We're looping over n's children
- n's estimate of the childrens' min is dropping

- Prizes A, B
- Note: a agent can be entirely rational without very representing or manipulating utilities and probabilities
- * $L = [p, A; (1-p), B]$

• Notation

- Preference: $A \succ B$
- Indifference: $A \sim B$

• Rational Preferences:

- Axiom of Transitivity: $(A \succ B) \wedge (B \succ C) \implies (A \succ C)$
- Orderability: $(A \succ B) \vee (B \succ A) \implies (A \sim C)$
- Continuity: $A \succ B \succ C \implies \exists p [p, A; 1-p, C] \sim B$
- Substitutability: $A \sim C \implies [p, A; 1-p, C] \sim [p, B; 1-p, C]$
- Monotonicity: $A \succ B \implies (p \geq q \iff [p, A; 1-p, B] \succeq [q, A; 1-p, B])$

- MEU(Maximum Expected Utility) Principle: Given any preferences satisfying these constraints, there exists a real-valued function U such that:

- Choose the action that maximizes expected utility
- Note: an agent can be entirely rational without very representing or manipulating utilities and probabilities
- Example: A lookup table for perfect tic-tac-toe, a reflex vacuum cleaner
- $U(A) \geq U(B) \iff A \succeq B$
- $U([p_1, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$

- i.e. values assigned by U preserve preference of both prizes and lotteries!

1.7 Markov Decision Processes

• Grid World

- A Maze-Like Problem where the agent lives in a grid and wall block the agent's path
- Noisy Movement: Actions do not always go as planned
- 80% of the time, the action North takes the action north if there is no wall
- 10% of the time, North goes West, 10% it will go east
- If there is a wall in the direction the agent would have moved, it will stay still
- Small living reward for each step
- Big reward comes at the end
- Goal: maximize sum of rewards

- A MDP is defined by

- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Probability that a from s leads to s' i.e. $P(s' \leftarrow s, a)$
 - Also called the model or the dynamics
- A reward function $R(s, a, s')$
- A start state
- Maybe a terminal state




Figure 1: Example $\alpha - \beta$ Pruning

- Who cares about n's value? MAX
- Let α be the best value that MAX can get at any choice point along the current path from the root.
- If n becomes worse than α , MAX will avoid it, so we can stop considering n's other children (it's already bad enough that this won't be played)

- α = max best option currently available, β = min best option currently available

- Formally with alpha-beta pruning:

```
def max-value(state, alpha, beta):
    initialize v = -inf
    for each successor of state:
        v = max(v, min-value(successor, alpha, beta))
        if v >= beta return v
    alpha = max(alpha, v)
    return v

def min-value(state):
    initialize v = +inf
    for each successor of state:
        v = min(v, max-value(successor, alpha, beta))
        if v <= alpha return v
    beta = min(beta, v)
    return v
```

- Alpha and Beta get passed down, only v gets passed up
- Pruning has no effect on the minimax value computed for the root
- With "perfect ordering", the time complexity drops to $O(b^{n/2})$ and it doubles solvable depth!

1.6 Searching with Expectimax and Utilities

- Idea: Uncertain outcomes controlled by chance, not an adversary. Sometimes the enemy doesn't play optimally and the penalty isn't too bad.

Remark 4. Up Triangle = Maximizer, Down Triangle = Minimizer, Circle = Chance Nodes

- In a deterministic single-agent search problem, we wanted an optimal plan, or sequence of actions from start to goal

- MDPs are non-deterministic search problems

- Markov generally means that given the present state, the future and past are independent

- For Markov decision processes, "Markov" means action outcomes depend only on the current state

- For MDPs, we want an optimal policy $\pi^*: S \rightarrow A$

- An optimal policy is one that maximizes utility when followed

- It's also reasonable to maximize the sum of the rewards

- However we can also prefer rewards now to rewards later

- One solution is to make the values of rewards decay exponentially

- Suppose we make the discounted reward less than one, then we can say reward at step 1 = 1, reward at step 2 = γ , reward at step 3 = γ^2

- Example: $\gamma = .5$, then reward for each future stage will decrease in half

- Example Problems

- Actions: East, West, Exit
- Q: For $\gamma = 1$, what is the optimal policy? 10, $\leftarrow, \leftarrow, \leftarrow, 1$
- This is because (10, 1), (10, 1), (10, 1), (10, 1), (10, 1)
- Q: For $\gamma = 0.1$, what is the optimal policy? 10, $\leftarrow, \leftarrow, 1$
- This is because (10, 0.0001), (1, 0.001), (0.1, 0.01), (0.01, 0.1), (0.001, 1)

- What if we have infinite utilities? We can use finite horizon which terminates episodes after a fixed T steps. This gives non stationary policies (π depends on time left) aka just declare an end.

- The value of a state s is: $V^*(s) = \max_a Q^*(s, a)$, expected utility starting in s and acting optimally

- The value(utility) of a q-state (s, a): $Q^*(s, a) = \sum_{s'} T(s, a, R(s, a, s')) + \gamma V^*(s')$ expected utility starting out having taken action a from state s and acting optimally. The reward of the action leading to s' plus it's discounted score

- The optimal policy: $\pi^*(s) = \text{optimal action from state s}$

- Propositional Logic is logic based on (and, or, not, iff, implies)

- Entailment: $(a = b$ is read as a entails b, or b follows from a) and it means every world where a is true, b is also true

- First order logic also includes (forall and exists)

1.8 Notes for Self

- When a heuristic cannot be defined correctly, and we need to fully explore a search space with all solutions on the lowest level, it stands to reason that A* will expand the same amount of nodes as DFS

- In the case of Pacman being in the same world as a single Box, both being at some coordinates (M,N), we can state that the state space would be $(MN)^2$ since pacman can be in MxN spots and the box can be in MxN spots

- The rules for a heuristics are the same as the rules for a measure

- In forward checking we only prune the domain that a assignment directly affects.